

Sapphire Local Runtime Benchmark

A reproducible local measurement, not a cluster-scale or GPU benchmark.

This manual describes the currently tested implementation in this repository and intentionally avoids claiming production-scale or distributed infrastructure.

It is a local runtime reference for learning, experimentation, scripting, and bounded automation work, not a verified multi-node training or cloud deployment platform.

Benchmark workload

The repository includes a local runtime benchmark that parses and interprets a 100-item loop repeatedly. One reference run on Python 3.12.10 produced result 4950 with median 0.21095 ms, minimum 0.2023 ms, and maximum 0.7817 ms.

How to reproduce

Run `python benchmarks/benchmark_runtime.py` from the repository root. Results vary with CPU load, Python version, operating system, and system state. This is intentionally a local benchmark, not a broad performance claim.

Boundary of claims

The project does not publish fabricated GPU, large-model, or multi-node throughput numbers. Those workloads are not executed or benchmarked by this repository.

Interpretation

A measured local loop time is informative for the current implementation, but not a guarantee of cross-language throughput or cluster-scale performance.

Current status summary

Sapphire version 1.0.5 is a Python-based local runtime and language with a working interpreter, bundled tools, and optional integrations. It is appropriate for learning, local automation, and experimental AI workflows. It is not presented as a distributed training stack, a production security-certified platform, or a globally managed service.

Results vary by machine, Python version, OS configuration, library availability, and workload. The documentation here reflects what is actually exercised in this repository and what remains a roadmap item.

Version 1.0.5. Reproduce local measurements with `benchmarks/benchmark_runtime.py`. Values depend on the machine and environment.