Sapphire Local Runtime Benchmark

A reproducible local measurement, not a cross-language, GPU, or cluster benchmark.

This manual describes the current local implementation. Feature names are not evidence of production scale, hardware availability, or successful external actions.

Workload

Parse and interpret a 100-item loop: let total = 0; for item in range(100) { total = total + item; } total;

Reference run

Python 3.12.10, 100 iterations, result 4950: median 0.21095 ms, minimum 0.2023 ms, maximum 0.7817 ms.

How to reproduce

Run python benchmarks/benchmark_runtime.py from the repository root. Results vary with Python version, operating system, CPU load, and hardware.

Excluded claims

The project does not publish fabricated MFU, tokens/second, 512-GPU throughput, or cross-language scores because those workloads are not executed by the repository benchmark.

Version 1.0.5. Reproduce local measurements with benchmarks/benchmark_runtime.py. Values depend on the machine and environment.